

Hierarchical Memory Orchestration with RAMSES for Robust Industrial AI Inference Systems

Geunsik Lim

Sungkyunkwan University

Suwon, South Korea

leemgs@g.skku.edu

Abstract—Serving models whose state spans GPU video RAM (VRAM), host dynamic RAM (DRAM), and NVMe is difficult because capacity pressure, asymmetric transfer bandwidth, queueing, imperfect prefetch, and synchronization interact under bursty load. We present RAMSES, a single-node hierarchical-memory orchestrator that combines GPU resource preallocation, NUMA-aware host allocation, and asynchronous block-aligned swapping. Its overlap-aware model separates resident demand from measured bidirectional transfer volume and uses two observable pressure ratios to select admission, prefetch, and eviction actions with hysteresis. On the reported two-A100 testbed, 16,000 request-level measurements across five models and four task types (five restarts each) reduce end-to-end inference latency by 34–41% relative to the PyTorch configuration; measured across the full memory hierarchy (GPU, CPU, and DRAM), whole-node energy per token falls by 47% and the energy–delay product by 70%, while token throughput rises by 61%. Model-load time and peak VRAM drop by 60% and 15%, respectively. On a named industrial defect-inspection task (MVTec AD) RAMSES matches baseline accuracy with output equivalence above 0.999. A 72-hour parameterized synthetic arrival trace yields fewer allocation failures and lower latency drift. These results are empirical QoS measurements for single-node serving, not field deployment, worst-case execution-time, control-stability, or functional-safety evidence. The model, measurement schema, and anonymized trace generator provide a reproducible protocol for studying when cross-tier coordination helps and when capacity or storage bandwidth remains the limiting resource.

Index Terms—Industrial AI, hierarchical memory, GPU offloading, NUMA, inference serving, workload trace

I. INTRODUCTION

MODEL serving becomes a memory-systems problem when parameters and key–value state exceed accelerator memory. Requests then couple GPU video RAM (VRAM), host dynamic RAM (DRAM), NVMe bandwidth, queueing, and synchronization. A policy that helps one configuration can hurt another because the active bottleneck changes with working-set demand, reuse, concurrency, and transfer direction.

This issue matters for industrial analytics such as inspection, predictive maintenance, and digital-twin support, where sustained availability and energy budgets are important. Our evidence, however, is a laboratory single-node evaluation driven partly by a parameterized synthetic arrival trace; it is not a factory field trial or a hard-real-time control-loop evaluation. We consequently report deadline misses only as Quality of Service (QoS), make no 2 ms or Time-Sensitive Networking compliance claim, and do not map request misses

to IEC 61508 Safety Integrity Levels.

Prior systems optimize paging, GPU–CPU offload, or request scheduling. RAMSES instead coordinates the VRAM–DRAM–NVMe path within one node while observing occupancy separately from actual bidirectional traffic. The design uses a lightweight state machine to choose admission, prefetch, and eviction actions rather than treating an analytical boundary as a guarantee.

The central research question is not whether another offload implementation can reduce memory use, but *when coordination across all three tiers changes end-to-end behavior and when physical limits make that coordination ineffective*. We answer it through a model–mechanism–evidence chain: the model identifies an operating regime, the controller changes a named set of actions, and the evaluation tests both the benefit and its predicted boundary.

The contributions are:

- *A falsifiable operating model.* The model separates resident demand from actual bidirectional traffic and accounts for overlap, queueing, reuse, and synchronization. Two dimensionless pressures predict a coordination-dominated region and capacity- and transfer-limited regions where RAMSES should provide little or negative benefit.
- *A regime-to-action controller.* RAMSES maps those observable pressures to admission, prefetch, and eviction actions using a specified 200 ms sampling interval, 5% hysteresis, and three-sample switching rule. This mapping, rather than its component mechanisms in isolation, is tested by the policy-off ablation.
- *A deployable single-node realization.* GPU preallocation, NUMA-local host allocation, and asynchronous block-aligned transfer are combined behind an LD_PRELOAD layer for unmodified PyTorch applications. We do not claim a hardened GPUDirect Storage or multi-node data path.
- *Boundary-seeking empirical evidence.* The evaluation crosses models and pressure regimes, includes a controller-disabled counterfactual and parameter sweeps, and reports absolute tails, uncertainty, failures, and synchronized whole-node (GPU, CPU, DRAM) energy. Reported maxima are configuration-specific effects, not guarantees.

The scope is deliberately narrow: RAMSES supports best-effort single-node inference. Multi-node fabrics, plant integra-

tion, certified energy measurement, real-time guarantees, and functional-safety qualification remain future work.

II. BACKGROUND AND RELATED WORK

The relevant prior work falls into three buckets, and we walk through them in that order: system-level optimizations, model-level optimizations, and the empirical observations from production-style profiling that ended up shaping RAMSES [1]–[4].

A. System-Level Optimization

On the system side, the targets are the obvious ones—hardware bottlenecks and Operating System (OS) overhead. In-memory file systems like `RAMFS` or `tmpfs`, paired with direct I/O, cut model load time [5], [6]. DeepSpeed and ZeRO-Offload spread parameters across devices to relieve VRAM pressure [7], [8]. Lower in the stack, OS-level knobs do real work: `jemalloc` heap allocation, slab allocators, `buffer_head` tuning—small things individually, but they reduce hidden allocation costs that profilers tend to miss [9]–[11]. Two more lines of work deserve a separate mention because they map onto what RAMSES eventually has to do: NUMA-aware scheduling [3], [12], and CPU–GPU offloading [13]–[16]. Both improve utilization once memory gets tight.

B. Model-Level Optimization

The model-side toolkit is, by now, well established. Quantization, pruning, LUT-based replacement: smaller compute footprint, smaller memory footprint [17], [18]. Mixed precision (FP16–FP32) lowers GPU memory pressure while keeping convergence quality intact [19]. Multi-GPU parallelism, FlashAttention, newer schedulers all help raw throughput [12], [20], [21]. Speculative execution and CPU-pooling cut redundant work and improve bandwidth utilization [22], [23].

C. Key Empirical Observations

When we profiled inference in setups close to production, a handful of inefficiencies kept turning up that current systems do not really address [2]–[4].

Cache sensitivity. Keeping in-memory caches warm reduced repeat-initialization time by something close to 70%. Flushing the cache wiped that out and forced a full reload [1], [5].

Synchronization overhead. Heavy slab-allocator activity together with frequent `futex()` calls produced allocation delays and lock contention under mixed workloads [9], [10].

NUMA and mapping latency. `mmap`-based parameter loading combined with cross-node DRAM access blew up latency the moment the working set ran past local memory [3].

Multiprocessing instability. Under DRAM pressure, aggressive multiprocessing fragmented memory and ran up the context-switch bill—throughput and stability both suffered [15].

Implications. Read together, these observations point to a single conclusion: VRAM, DRAM, synchronization, and I/O are tightly coupled, and none can be tuned in isolation without perturbing the others. For industrial CPS, the same instability propagates outward into control loops, where it stops being a performance issue and becomes a reliability

TABLE I
STRUCTURAL COMPARISON BETWEEN RAMSES AND REPRESENTATIVE SYSTEMS. RAMSES UNIQUELY ENABLES FULL COORDINATION ACROSS VRAM, DRAM, AND NVME, WHILE PRIOR SYSTEMS PROVIDE ONLY PARTIAL SUPPORT. ABBREVIATIONS: PART. (PARTIAL SUPPORT), ARCH. (ARCHITECTURAL-LEVEL ANALYSIS ONLY), N/A (NO EMPIRICAL EVALUATION DURATION REPORTED), AND TP/W (THROUGHPUT PER WATT).

System	Tier Scope	NUMA aware	Regime Model	Eval. Dur.	Energy Metric
FlexGen	GPU–CPU–disk	×	None	Short	×
SwapAdvisor	GPU–CPU	×	None	Short	×
NEO	Partial	×	None	Short	×
SpecOffload	GPU–NVMe	×	None	Short	×
vLLM	GPU-only	×	None	Short	Part.
Ind. Edge/DT	Arch.	Part.	None	N/A	×
RAMSES	Full	✓	Analytical	24–72h	TP/W

issue. Coordinated VRAM–DRAM–NVMe management with explicit regime awareness was, at that point, the obvious next step.

D. Limitations of Prior Approaches

Most prior work tackles one layer or one resource at a time. That makes compound bottlenecks awkward to fix—the obvious examples are GPU OOM errors, memory-mapping stalls, and NUMA-induced delays [1], [11], [15]. The picture gets worse in multi-tenant or multi-model deployments [3], [12], [24]. Memory-management and offloading systems—FlexGen, SwapAdvisor, SpecOffload, eLLM—do orchestrate GPU–CPU–NVMe transfers, but reactively, and without a formal regime characterization [1], [25], [26]. System-level scheduling improves placement, but again, no unified regime picture [8], [11], [27], [28]. Filling that gap is what we set out to do: regime-aware analytical conditions plus a unified VRAM–DRAM–NVMe orchestrator.

E. Positioning of RAMSES

Table I lays out the structural differences between RAMSES and a set of representative systems along five axes: tier scope, NUMA awareness, the formality of regime modeling, evaluation duration (Eval. Dur.), and energy reporting (throughput-per-watt, TP/W). “Part.” marks partial support; “Arch.” marks architectural-level analysis only. The general picture is that prior LLM-serving frameworks tune for throughput or stretch memory via partial GPU–CPU coordination—but none of them evaluates the same combination of intra-node NUMA placement, resource preallocation, and pressure-driven policy switching.

Several recent TII papers cover complementary parts of the stack without touching the intra-node memory dynamics that RAMSES targets. Tao *et al.* [29] model latency at the digital-twin synchronization layer, but they do not engage GPU VRAM–DRAM swap dynamics or single-node pressure-driven inference. Cao *et al.* [30] survey cloud–edge resource coordination and explicitly bracket out the intra-node memory hierarchy; RAMSES supplies the missing memory-subsystem layer. Chen *et al.* [31] use deep RL to allocate mobile-edge compute at the task-scheduling level. Inside each node, RAMSES is what manages VRAM/DRAM/NVMe once the task has been placed. Wu *et al.* [32] study accuracy–latency

tradeoffs through model splitting, with each device treated as a black box. RAMSES empirically reduces the per-device delay that such split-inference schedulers implicitly assume. **Lu et al. [33]** reduce communication overhead in federated learning for digital twins; on each participating node, RAMSES reduces local inference latency and VRAM pressure, which shortens per-round computation. **Chekired et al. [34]** propose hierarchical fog-layer scheduling for smart-factory IoT data and treat inference nodes as latency constants—which is exactly where RAMSES plugs in by giving a regime-aware way to derive those constants from VRAM–NVM behavior. **Deng et al. [35]** use RL for dynamic edge-resource allocation, modeling inference cost at the request-dispatch level; RAMSES supplies measured per-node latency distributions that an RL policy could use as observations.

The pattern is consistent: prior TII work coordinates at the network, task, or device level and treats inference nodes as opaque units with fixed latency. RAMSES sits one layer below all of this, characterizing regime-dependent behavior across VRAM–DRAM–NVM inside a single serving node and producing empirical latency distributions that upstream schedulers can use.

F. Distinct Contributions of RAMSES

What is distinctive in RAMSES is the combination of two things that usually appear separately. First, hierarchical memory orchestration. Second, an analytical regime characterization built on a latency decomposition over compute, locality, transfer, and synchronization. Going past conventional serving optimizations, we identify observable pressure ratios that classify capacity and transfer stress. The three internal modules play distinct roles—the GPU Process Allocator handles context conflicts, the Task Memory Manager handles NUMA-aware allocation, and the GPU Booster handles swap-aware asynchronous control—but each is judged against the same regime boundaries. Where systems-venue work tends to optimize for throughput or for one specific layer, the focus here is single-node QoS and sustained-load behavior under a parameterized synthetic workload.

III. RAMSES: ARCHITECTURE AND KEY COMPONENTS

LLM inference relies on tight cooperation across a heterogeneous set of resources: GPUs, VRAM, DRAM, NVMe, and Central Processing Units (CPUs). Existing serving stacks and runtimes do not handle resource contention as a first-class concern, which is why we keep seeing the same failure modes: Compute Unified Device Architecture (CUDA) out-of-memory (OOM) errors, context-switch stalls, and memory-mapping overhead. Earlier attempts to plug these gaps include GPU–CPU offloading [26], speculative offloading [14], and cache-centric management [13]. Useful as those are, most stop at single-tier transfers (typically VRAM–CPU) or address only narrow scenarios such as key-value (KV) cache management.

We propose **RAMSES** as a way past these limits. The design borrows the OS notion of memory ballooning and




Fig. 1. Industrial hierarchical memory architecture of RAMSES for laboratory workloads motivated by industrial analytics. The data plane (solid arrows) coordinates GPU VRAM, host DRAM, and NVMe for three industrial workloads: (1) digital twin state synchronization aligned with PLC scan cycles, (2) predictive maintenance triggered by anomaly bursts from IIoT sensors, and (3) real-time visual inspection under bounded inference deadlines. The control plane (dashed arrows) includes a Regime Analyzer estimating (α, β) and a Policy Engine enforcing tier-aware placement, prefetch, and swap decisions. Three modules handle the actual orchestration: the *GPU Process Allocator* pre-reserves CUDA context to remove context-conflict stalls, the *Task Memory Manager* performs NUMA-aware allocation with predictive DRAM offloading, and the *GPU Booster* drives block-aligned asynchronous VRAM–NVM swap. The diagram supplies workload motivation only; it does not depict a deployed control loop, TSN integration, or a validated hard deadline.

applies it to LLM serving. What separates RAMSES from prior work is that it manages the whole VRAM–DRAM–NVMe hierarchy together: GPU context conflicts, NUMA-induced latency, and inefficient swapping are addressed in the same control loop instead of patched independently. RAMSES also reallocates GPU and system memory on the fly, while staying lightweight enough to avoid the deployment burden of full distributed-processing frameworks. Three modules sit at the core of the design—the **GPU Process Allocator**, the **Task Memory Manager**, and the **GPU Booster with swap support**—and they are illustrated in Fig. 1.

A. GPU Process Allocator

In multi-model serving, two bottlenecks dominate: **GPU context loading conflicts** and **VRAM interference**. CUDA drivers do context initialization and memory mapping with *lazy binding*, which is fine for a single workload but becomes painful under concurrent execution—startup delays grow, and they show up directly in user-facing latency. NEO [13] addresses cache management but leaves the GPU-level context conflict question open. The **GPU Process Allocator** is our answer to this. Rather than waiting for runtime to resolve these conflicts, it reserves GPU execution resources up front: context IDs, memory pools, and kernel queues are pre-assigned. It also keeps memory segments separated across models, which keeps cache pollution down, and it co-schedules I/O-intensive and compute-intensive models so the GPU is not idle waiting on either side.

In practice, the allocator runs three policies. **Pre-defined execution slot reservation** hands each model a context ID, memory pool, and kernel queue ahead of time, so CUDA initialization no longer happens on the critical path. **Context**

¹The anonymized review artifact bundles the LD_PRELOAD orchestrator layer, the synthetic-trace generator, baseline invocation scripts, the request-level measurement schema, and the analysis programs used to produce the reported results.

Algorithm 1 GPU Process Allocator Pseudocode**Require:** Set of models $\mathcal{M}$, GPU resource pool $\mathcal{R}$

- 1: Initialize free context IDs, memory pools, and kernel queues
- 2: **for** each model $m \in \mathcal{M}$ **do**
- 3: Reserve a context ID, memory pool, and kernel queue for m
- 4: Assign a priority based on whether m is I/O-bound or compute-bound
- 5: **while** tasks remain **do**
- 6: Dequeue the next task t based on priority
- 7: Launch t on its reserved context
- 8: **if** t finishes **then**
- 9: Return its resources to $\mathcal{R}$

separation with fixed memory pools assigns models to separate configured regions; the finite experiments show lower fragmentation and cross-model cache interference. **Priority-based scheduling** orders execution between I/O-bound and compute-bound models so neither stalls the other and idle GPU cycles shrink. The result is a small extension of process pinning that helps preserve **Quality of Service (QoS) stability** and gives more **predictable latency** when several models share one GPU [8]. Algorithm 1 summarizes the end-to-end pre-reservation and dispatch flow.

B. Task Memory Manager

Memory usage during LLM inference shifts back and forth between DRAM and VRAM as input and generation lengths change. eLLM [26] responded to this by extending the CPU buffer; PIE [22] took a speculative-swap approach. Neither, though, paid attention to NUMA-induced latency or to optimization within DRAM itself.

Our **Task Memory Manager** closes that gap. It pairs task-level memory control with NUMA-aware policies: by aligning processes with the right memory nodes, remote-access delay is reduced, and a jemalloc-based custom heap separates small objects from large tensors so that fragmentation stays manageable [9], [10]. On top of this, the module looks ahead: tensors with low expected access probability are offloaded preemptively, and a small predictive model of memory pressure feeds anticipatory decisions [18]. We model memory access latency as

$$T_{\text{access}} = T_{\text{local}} + \gamma \cdot T_{\text{remote}},$$

where γ is the penalty factor from a remote NUMA access. Bad allocation policy here cascades—excessive *page reclaim*, *mmap* failures, unnecessary *buffer eviction*. Existing fixes mostly fall back to reactive GPU offloading or memory compression once VRAM gets tight; they leave DRAM behavior and NUMA locality on the table.

The Task Memory Manager has four pieces. **NUMA-aware memory mapping** pins process and thread data to local NUMA nodes through system calls such as `mbind()` and `numa_alloc_onnode()` so that remote-memory traffic is minimized. **jemalloc-based heap optimization** tunes pool size, allocation granularity, and cache-aware placement, with small objects kept apart from large tensors. **Task-level redistribution** watches per-request memory consumption and of-

Algorithm 2 Task Memory Manager Pseudocode**Require:** Execution graph G , memory threshold τ

- 1: Bind threads to local NUMA nodes using OS calls (e.g., `numa_alloc_onnode`)
- 2: Initialize separate jemalloc heaps for small objects and large tensors
- 3: **for** each incoming request r **do**
- 4: Monitor VRAM usage and fragmentation
- 5: **if** usage $> \tau$ **then**
- 6: Identify tensors in G with low immediate reuse probability
- 7: Offload these tensors to DRAM/NVMe
- 8: Use a lightweight predictor to anticipate memory pressure based on current utilization and I/O depth
- 9: Pre-allocate VRAM/DRAM resources for the next request

floods non-critical tensors to DRAM before VRAM saturates. A **lightweight memory pressure prediction model** reads live utilization and I/O queue depth to predict shortages and trigger mitigation early. Together, these free up VRAM and—a point worth highlighting—make DRAM–VRAM use more efficient as a unit, with less fragmentation and fewer *page reclaim* events. RAMSES is therefore able to keep memory provisioning stable and access latency low even when several models and processes contend for the same node. Algorithm 2 states the full NUMA-aware allocation and predictive offloading procedure.

C. GPU Booster with Swap Support

VRAM capacity remains one of the hardest constraints in serving large models. FlexGen [25] and SwapAdvisor [1] support GPU–CPU swapping, but neither is tuned to GPU memory hierarchies in particular.

The **GPU Booster with swap support** targets this bottleneck by selecting inactive tensors and scheduling block-aligned asynchronous transfers. It is implemented as an `LD_PRELOAD` interception layer over the CUDA runtime and caching allocator: allocation, free, and host–device copy calls are intercepted so that tensor placement and movement are managed without modifying the model or the serving framework. Intercepted entry points, buffer-registration details, and exact driver/filesystem versions are documented with the review artifact.

Conventional offloading typically pays a noticeable transfer cost, and OS-level swapping is unaware of GPU memory semantics, which limits how well it can do. The GPU Booster instead uses execution-graph reuse distance to rank eviction candidates, dedicated asynchronous CUDA streams to overlap transfer with compute, and 4 MB block alignment matched to the NVMe erase-block granularity; the block-size choice is examined in the sensitivity study of Section IV. On hosts that expose GPUDirect Storage, tensors move over a direct VRAM–NVMe DMA path; otherwise the layer falls back to a pinned-host-buffer staged path with the same asynchronous overlap. Partial reload fetches only the portion of a swapped tensor required by the next kernel, and output equivalence against the unoptimized baseline is checked bitwise for FP32

TABLE II

COMPARISON OF RAMSES WITH EXISTING BASELINES. EACH ROW IDENTIFIES A BASELINE SERVING STACK, ITS KEY OPTIMIZATION TECHNIQUE (*Key Features*), THE REASON IT WAS SELECTED FOR EVALUATION (*Rationale*), AND WHICH RAMSES MODULE ADDRESSES THE CORRESPONDING BOTTLENECK (*Targeted by RAMSES*). CITATIONS REFER TO THE BASELINE SYSTEMS.

Method	Key Features	Rationale	Targeted by RAMSES
Default (PyTorch)	No optimization	Ground truth reference	Overall performance gain
FlexGen [25] (ICML'23)	CPU offloading	VRAM expansion via offload	GPU Booster (swap)
SwapAdvisor [1] (ASPLOS'20)	GPU-CPU swap	Canonical swap optimization	Task Memory Manager
NEO [13] (MLSys'25)	KV cache offload + sched.	Recent tier utilization	Hierarchical (full) tier use
SpecOffload [14] (arXiv'25)	Cross-tier swap min.	Latest transfer optimization	NVMe-VRAM transfer path
vLLM [24] (SOSP'23)	KV cache + streaming	SOTA inference optimization	End-to-end latency

Algorithm 3 GPU Booster with Swap Support Pseudocode

Require: Execution graph G

- 1: **for** each tensor t in G **do**
- 2: **if** t is not required in upcoming compute steps **then**
- 3: Align t to 4MB boundaries
- 4: Initiate asynchronous DMA transfer of t to NVMe storage
- 5: Mark t as swapped out
- 6: **while** serving requests **do**
- 7: **if** the next computation step reaches a swapped-out tensor t **then**
- 8: Partially reload t via asynchronous DMA from NVMe
- 9: Resume computation once the needed portion of t has been fetched




Fig. 2. Operating map for the two pressure ratios. Here $\alpha = D/C_f$ denotes capacity pressure and β is the ratio of non-overlapped transfer time to compute, memory, and synchronization time. The dashed lines are operational classification thresholds used by the controller, not feasibility, determinism, or stability guarantees. Under the evaluated workloads RAMSES operates in the coordination-dominated region ($0.6 < \beta < 0.9$), whereas the baselines run at higher mean transfer pressure and approach the transfer-limited boundary ($\beta \rightarrow 1$) under burst load; the corresponding measured behavior is reported in Section IV.

and within the precision tolerance for FP16.

The goal is to serve models whose state exceeds VRAM capacity while keeping exposed transfer latency low. Algorithm 3 summarizes the control flow.

D. Overlap-Aware Operating Model

We use the following model as an engineering classification tool, not as a real-time guarantee. Let $D(t)$ be the bytes resident or demanded by a step, $C_f = C_{\text{VRAM}} + C_{\text{DRAM}}$ the

usable fast-tier capacity, $V_{\uparrow}(t)$ and $V_{\downarrow}(t)$ the measured read and write transfer volumes, $h(t)$ the prefetch-hit fraction, and $q(t)$ the measured storage queueing delay. These quantities are independent observations; in particular, occupancy is not treated as bytes transferred. The capacity pressure is

$$\alpha(t) = D(t)/C_f, \quad (1)$$

and the transfer pressure is

$$\beta(t) = \frac{L_0 + q(t) + V_{\uparrow}(t)/B_{\uparrow} + V_{\downarrow}(t)/B_{\downarrow}}{T_{\text{comp}}(t) + T_{\text{mem}}(t) + T_{\text{sync}}(t)}. \quad (2)$$

Here L_0 is fixed request/setup latency and $B_{\uparrow}$ and $B_{\downarrow}$ are effective, direction-specific bandwidths measured with the same block size and queue depth as the experiment. All symbols are observables; no critical value is defined by substituting a boundary back into itself.

Asynchronous transfer overlaps computation. Accordingly, predicted service time is

$$\hat{T}_{\text{total}} = T_{\text{sync}} + T_{\text{mem}} + \max\{T_{\text{comp}}, (1-h)T_{\text{xfer}}\}, \quad (3)$$

$$T_{\text{xfer}} = L_0 + q + V_{\uparrow}/B_{\uparrow} + V_{\downarrow}/B_{\downarrow}. \quad (4)$$

The max operator prevents double counting overlapped DMA and compute. Block rounding is included in the measured volumes. We call observations with $\alpha > 1$ capacity-limited and observations with $\beta \geq 1$ transfer-limited; the remaining points are coordination-dominated. These are operational labels, not theorems. They neither imply zero swapping nor a worst-case bound. Fig. 2 summarizes the resulting classification. Section IV reports the operating regime the evaluated systems occupy; per-request measured (α, β) values and the residual of Eq. (3) against end-to-end latency are logged by the instrumentation of Section IV-B and released with the artifact.

E. Online Policy and Complexity

Every 200 ms the analyzer reads allocated bytes, transfer counters, queue depth, effective bandwidth, and exponentially weighted compute, memory, and synchronization times. It computes Eqs. (1)–(2) in constant time. To avoid oscillation, a state change requires three consecutive samples beyond a 5% hysteresis band. In the capacity-limited state the policy rejects admissions that exceed the configured memory reserve and evicts the lowest predicted-reuse tensors. In the transfer-limited state it stops speculative eviction, reduces prefetch depth, and prioritizes reads. In the coordination-dominated state it enables asynchronous prefetch and NUMA-local placement. The objective is lexicographic: prevent allocation failure, then minimize measured p99 latency, then transfer bytes.

The reuse score is an exponentially weighted count of accesses over the last 32 graph steps; it has no learned parameters or offline training. Candidate maintenance costs $O(\log n)$ per tensor access for a heap of n evictable tensors; sampling, regime classification, and a policy decision are $O(1)$, and metadata occupies $O(n)$. To isolate the controller from its mechanisms, we evaluate a policy-off configuration that keeps all three data-plane modules and fixed placement parameters but disables regime switching; its trace and seed are shared with the full-system run. Disabling the policy forfeits the coordination-region behavior, so the system drifts toward the transfer-limited boundary under burst load, consistent with the sustained-load results in Section IV; the paired per-run records are included in the artifact. The measured 1.68 ms is mean controller work per serving request and is not compatible with a 2 ms end-to-end claim; we therefore make no such claim.

F. Scope of Timing and Safety Claims

The measured latency distributions are QoS observations for non-safety AI serving. They are not WCET bounds, deterministic-latency guarantees, control-loop stability proofs, or evidence of IEC 61508 Safety Integrity Level (SIL) compliance. A deadline-miss fraction is not a probability of dangerous failure on demand. Establishing a safety claim would additionally require a defined safety function and demand mode, hazard analysis, diagnostic coverage, common-cause and safe-state assumptions, a plant-specific delay margin, and hardware- or software-in-the-loop validation. RAMSES must therefore be kept outside a safety function unless independently qualified by the integrator.

G. Ablation Analysis

The three modules of RAMSES contribute on their own, but they also reinforce each other when combined. To show this, we ran an **ablation study** over four configurations: full RAMSES, RAMSES without the GPU Process Allocator, without the Task Memory Manager, and without the GPU Booster. Table III summarizes the outcome—every ablation degrades performance, and the breakdown matches the design intent of each module. Removing the **GPU Process Allocator**, in particular, has the largest impact on inference latency ($35.0\% \rightarrow 20.5\%$, a 14.5 percentage-point drop), consistent with its role of pre-resolving CUDA context conflicts on the critical path. Removing the **Task Memory Manager** has the largest impact on VRAM reduction ($15.0\% \rightarrow 5.0\%$, a 10.0 percentage-point drop), consistent with its NUMA-aware allocation and predictive DRAM offloading. Disabling the **GPU Booster** produces an intermediate degradation across both latency and VRAM, consistent with its swap-aware role between VRAM and NVMe. The pattern lines up cleanly with the design intent: each of the three modules dominates the metric it was built to target [18].

IV. EVALUATION

This section reports a quantitative evaluation of **RAMSES** relative to existing approaches. Every experiment is run five times under the same hardware setup. Tables report absolute request-level percentiles and run-level means with 95% confidence

TABLE III
ABLATION STUDY RESULTS: IMPACT OF DISABLING RAMSES MODULES

Configuration	Load Time Reduction	Latency Reduction	VRAM Reduction
Full RAMSES	58.7±1.8%	35.0±1.2%	15.0±0.9%
Without GPU Allocator	40.1±2.1%	20.5±1.6%	14.1±1.0%
Without Memory Manager	45.0±1.9%	28.0±1.4%	5.0±0.7%
Without GPU Booster	49.5±2.0%	25.4±1.5%	10.2±0.8%

Values are mean ± standard deviation over five runs. Raw per-run observations and the analysis scripts are provided in the review artifact.

intervals; standard deviations remain available in the artifact. All request-level latency, whole-node energy, operating-model, controller-ablation, and per-model results are computed by the released scripts from `code/data/actual/raw.jsonl`—16,000 request-level records spanning eight serving configurations, five models, four task types, and five restarts—so each such value traces back to primary measurements rather than to normalized plots. The named-industrial-task accuracy, the parameter-sensitivity sweep, the model-load time and peak-VRAM measurements (Table VI), and the 72-hour arrival-trace QoS series (Table VIII) are likewise released as measured CSVs under `code/data/actual/`. Raw per-run observations, the request-level logs, and the analysis scripts that compute confidence intervals and significance tests are provided in the review artifact so that every summary in this section can be recomputed from primary data.

A. Research Questions and Claim–Evidence Alignment

We organize the evaluation around four questions tied directly to the contributions in Section I:

- RQ1: **Deadlines and pressure** ratios distinguish the three regimes, and how accurately does the latency model predict held-out request latency? We report operating-point coverage, MAE/RMSE, and residuals.
- RQ2: **Values of measured state switching** add value beyond the same preallocation, NUMA, and transfer mechanisms at fixed settings? The same-trace `ramses_policy_off` run is the counterfactual, and parameter sweeps test dependence on a single chosen threshold.
- RQ3: **Where the boundary** RAMSES improve latency, memory, failures, and whole-node energy, and where does its overhead cease to pay? We retain non-gains and report absolute P50–P99.9 latency and failure denominators.
- RQ4: **Does model provenance** serve model outputs and named-task accuracy? Baseline/RAMSES outputs are paired on identical inputs; no accuracy value is inferred from a latency experiment.

Synthetic values under `code/data/expected/` are pre-experiment hypotheses used only to size the matrix and exercise the pipeline. They are excluded from every paper table, figure, statistical test, and claim. A result enters this section only when request-level measurement provenance exists under `code/data/actual/`; unsupported configurations remain explicitly missing rather than being filled from an expectation. **Testbed.** Unless stated otherwise, the platform is a laboratory server, not a deployed factory-edge system: two NVIDIA

TABLE IV

ABSOLUTE PER-TASK LATENCY PERCENTILES (TTFT) ON THE PRINCIPAL LLAMA-4 17B WORKLOAD, FROM RAW REQUEST-LEVEL LOGS. *Means across $n = 5$ runs (95% CI). HW: $2 \times$ NVIDIA A100 (80GB), 512GB DRAM, NVMe Gen4; SW: PyTorch 2.4, CUDA 12.2, vLLM 0.5.3.*

System	Median (ms)	P95 (ms)	P99 (ms)	P99.9 (ms)	Max (ms)
PyTorch (Default)	871.6	980.4	1003.7	1011.5	1012.4
FlexGen	792.8	914.0	957.8	979.7	982.2
SwapAdvisor	694.9	784.7	823.7	837.4	838.9
NEO	745.1	838.3	921.8	942.9	945.2
SpecOffload	668.0	775.4	792.3	833.8	838.4
vLLM	590.0	653.7	696.9	706.3	707.4
RAMSES	563.0	624.9	671.4	686.5	688.2
RAMSES (policy-off)	649.1	734.6	824.8	825.1	825.1

A100 (80GB) GPUs, an Intel Xeon Gold 6440 CPU (40 cores), 512 GB ECC DRAM, and a Gen4 NVMe SSD running PyTorch 2.4, CUDA 12.2, and vLLM 0.5.3. Because vLLM 0.5.3 predates native Llama-4 support, we applied a compatibility patch that registers the Llama-4 model class and adapts the configuration and weight loaders; the patch and the exact invocation are included in the review artifact so the comparisons are reproducible. The principal workload is Llama-4 (Maverick-17B-128E-Instruct) [36]. The machine has redundant power, ECC protection, and sustained-load thermal control, which is what 24/7 industrial operation actually looks like. Single-GPU ablations are flagged explicitly as “ $1 \times$ A100”.

For comparison we use the standards—FlexGen [25] and SwapAdvisor [1]—together with newer entries: NEO [13], SpecOffload [14], and vLLM [24]. Table II captures the comparison scope.

B. Request-Level Measurement Methodology

To prevent incompatible latency scales from being pooled, each request-level record is assigned to exactly one task: *scoring* (one forward pass, no generation), *continuation* (one warm decode step), *TTFT* (request arrival through first emitted token), or *generation* (arrival through the final token). For every task and system we log the model, numerical precision, input/output token counts, batch size, concurrency, request count, and run identifier, following the schema in `./code/measurement-schema.json`. From these logs the released analyzer computes absolute median, P95, P99, P99.9, and maximum latency per task and system directly from timestamps; no value is derived from normalized bars. Each comparison uses five independently restarted runs and reports both cold- and warm-cache protocols, and failed requests are retained in the failure-rate denominator. Table IV reports the absolute percentiles for the principal workload, and Table V reports the measured pressure ratios and the prediction error of the operating model against end-to-end latency.

RQ1: predictive validity. Table V evaluates the operating model directly. The overlap-aware prediction of Eq. (3) tracks measured end-to-end latency with a mean absolute error of 5.1 ms (RMSE 5.9 ms) for RAMSES and at most 8.9 ms (RMSE 10.1 ms) across every system—below 1.1% of the corresponding median latencies in Table IV. The measured transfer pressure orders the systems as the model predicts:

TABLE V

OPERATING-MODEL VALIDATION: MEASURED MEAN PRESSURE RATIOS, PREDICTION ERROR OF EQ. (3) AGAINST END-TO-END LATENCY, AND PREFETCH HIT RATE. *Means across $n = 5$ runs (95% CI). HW: $2 \times$ NVIDIA A100 (80GB), 512GB DRAM, NVMe Gen4; SW: PyTorch 2.4, CUDA 12.2, vLLM 0.5.3.*

System	α	β	MAE (ms)	RMSE (ms)	Prefetch hit rate
PyTorch (Default)	0.78	0.80	8.9	10.1	0.78
FlexGen	0.76	0.73	8.1	9.0	0.82
SwapAdvisor	0.78	0.77	6.5	7.4	0.83
NEO	0.74	0.74	6.7	7.9	0.81
SpecOffload	0.75	0.72	5.6	6.5	0.86
vLLM	0.78	0.72	4.8	5.8	0.84
RAMSES	0.74	0.66	5.1	5.9	0.96
RAMSES (policy-off)	0.77	0.69	6.0	6.9	0.93




Fig. 3. Measured operating points on the pressure map. Each marker is a measured (α, β) pair per serving configuration on the principal workload; dashed lines are the classification thresholds. RAMSES operates in the coordination-dominated region, and Table V reports the prediction error of Eq. (3) against measured latency.

RAMSES runs at $\beta = 0.66$, inside the coordination-dominated band ($0.6 < \beta < 0.9$), whereas the baselines sit at higher mean transfer pressure ($\beta = 0.72$ – 0.80) and cross into transfer-limited behavior during bursts, and RAMSES attains the highest prefetch-hit rate (0.96). These residuals and per-request (α, β) values are recomputed from the released logs.

C. Model, Policy, and Energy Instrumentation

Each request-level record additionally carries measured α , β , predicted latency, direction-specific effective bandwidth, storage queueing, prefetch hits/misses, and bytes read/written. From these the analyzer computes the residual of Eq. (3) against measured latency, its mean absolute error and root-mean-square error, and the prefetch-hit rate. The policy-off configuration retains all three data-plane modules and fixed placement parameters while disabling only regime switching, and each paired run reuses the same request trace and seed as its full-system counterpart.

Energy is integrated on the same monotonic clock as the request telemetry. The primary energy results are *synchronized whole-node* measurements: GPU energy from NVML and CPU-package and DRAM energy from Intel RAPL are summed over the same window by `code/collect_energy.py`, so activity that RAMSES shifts from the GPU to the host is captured rather than hidden. Each run records the instrument model, sampling

rate, clock alignment, idle-subtraction procedure, calibration uncertainty, and dropped-sample count. The measurement schema, collector, and analysis programs are included under `code/`, and the raw per-request energy is in `code/data/actual/raw.jsonl`. NVMe rail power and certified rack-level absolute accuracy (IEC 62053) remain future work.

D. Experimental probes across memory-dominated regimes

Four metrics drive the evaluation: (1) **model loading time**; (2) **inference latency**, captured as Time-To-First-Token (TTFT) and sustained throughput; (3) **memory utilization efficiency**, expressed as VRAM/DRAM occupancy and fragmentation, where fragmentation = $1 - \text{largest_free_block}/\text{total_free}$; and (4) **system stability**, i.e., VRAM overflows and GPU restarts under sustained load.

E. Baseline selection across heterogeneous orchestration regimes

The evaluated baselines are memory-centric and cross-tier serving systems: the PyTorch default, PagedAttention/vLLM [24], FlexGen [25], SwapAdvisor [1], and the more recent cross-tier offloading methods NEO [13] and SpecOffload [14]. With RAMSES and its policy-off variant these form the eight serving configurations reported throughout this section. Orca’s continuous-batching scheduler [21] is discussed as related work but is not benchmarked here, because it targets request scheduling rather than the VRAM–DRAM–NVMe memory path that RAMSES addresses. Selection criteria favor architectural diversity and direct relevance to regime-aware coordination.




Fig. 4. Consolidated experimental results across baselines and RAMSES on Llama-4 17B. All bars are normalized to the PyTorch *Default* baseline (=100%), where (a) load time, (b) peak VRAM usage, and (c) end-to-end inference latency are reported as percentages—lower is better. (d) Failures are absolute counts of VRAM overflows and GPU restarts observed during a 72-hour sustained run (numerical labels above the bars). The RAMSES column (rightmost, highlighted) achieves the lowest load time, latency, and failure count across all configurations.

F. Regime-dependent performance characteristics

Below we walk through RAMSES across the analytically defined operating regimes, showing how regime-aware coordination translates into performance, stability, and reliability gains.

TABLE VI

MEASURED MODEL-LOAD TIME AND PEAK VRAM PER SYSTEM (MEAN OVER 2,000 REQUESTS EACH) ON THE LLAMA-4 17B WORKLOAD. RAMSES REDUCES MEAN LOAD TIME BY 59.9% AND PEAK VRAM BY 15.1% RELATIVE TO PYTORCH. Means across $n = 5$ runs (95% CI). HW: $2 \times$ NVIDIA A100 (80GB), 512GB DRAM, NVMe Gen4; SW: PyTorch 2.4, CUDA 12.2, vLLM 0.5.3.

System	Load-time (s)	Peak VRAM (GB)
PyTorch (Default)	38.5	29.2
FlexGen	35.1	27.1
SwapAdvisor	33.2	26.9
NEO	31.7	26.6
SpecOffload	30.1	26.3
vLLM	27.8	26.0
RAMSES	15.4	24.8
RAMSES (policy-off)	22.0	25.7

1) Capacity-limited regime: model initialization under memory pressure

On Llama-4 17B, RAMSES cuts mean model-load time by 59.9% relative to PyTorch (38.5 s to 15.4 s; Table VI), averaged over 2,000 requests. Against FlexGen, SwapAdvisor, NEO, SpecOffload, and vLLM the reduction is 56.0%, 53.4%, 51.2%, 48.7%, and 44.4%, respectively. The source of these gains is twofold: CUDA context initialization no longer happens on the critical path, and NVMe–VRAM transfers are tuned.

2) Locality and fragmentation analysis

Fig. 4(b) shows that the *jemalloc*-based DRAM optimization in RAMSES cuts host-side heap fragmentation and allocation churn. On the VRAM side, the gains come from CUDA allocator tuning combined with selective GPU swapping, lowering mean peak VRAM by 15.1% (29.2 GB to 24.8 GB; Table VI). Inference failure rate stays under 0.5% across configurations.

3) Coordination-dominated regime: latency reduction beyond raw capacity

On the principal Llama-4 17B TTFT workload, RAMSES cuts median latency by 35.4% (871.6 ms to 563.0 ms) and raises token throughput by 61% relative to PyTorch (Table IV; all records at batch 1). Against the competing serving stacks the median-latency reduction ranges from about 5% over vLLM to 29% over FlexGen, confirming that cross-tier coordination reshapes end-to-end latency rather than merely freeing capacity.

4) Stability collapse and recovery across regime boundaries

Fig. 4(d) summarizes a continuous 72-hour stress test. RAMSES shows zero VRAM overflows and no inference crashes, whereas baseline systems hit 5–12 failures. The full 72-hour time series is released with the artifact; we report the failure counts as a reliability indicator and do not claim statistical stationarity or a causal mechanism from these counts alone.

5) Long-Duration QoS under a Parameterized Synthetic Trace

The 72-hour run uses the parameterized synthetic generator described below. It is not a PLC, TSN schedule, plant, or sensor-stream evaluation. We use one consistent QoS event: a request whose end-to-end latency exceeds $1.5\times$ the median of its fixed workload configuration. Over 4,320 one-minute




Fig. 5. Energy–latency tradeoff under varying batch sizes. Each gray marker is one (energy, latency) operating point from a baseline serving stack (FlexGen, SwapAdvisor, NEO, SpecOffload, vLLM, and the PyTorch *Default*) measured across the batch-size sweep, and the green curve is the convex hull of RAMSES operating points over the same sweep. RAMSES dominates the baseline cloud in both axes simultaneously—lower latency at equal or lower energy per inference—shifting the achievable energy–latency frontier inward.

records per system, RAMSES records a 0.60% miss fraction (26 violations) versus 8.70% (376 violations) for the PyTorch reference. We report this only as an empirical request-level QoS metric; it is unrelated to functional-safety failure probability. Mean p99 latency drift is 0.97% versus 5.07% for the reference, with no observed CUDA restart in this finite run (Table VIII).

G. Multi-GPU Scope

The contribution and all primary comparisons are scoped to one serving node. Two-GPU measurements use tensor placement within that node. We do not claim 4/8-GPU or fleet-level generality because the available data do not include fabric saturation, link or device failure injection, recovery time, or communication-tail analysis.

H. Regime invariance across model architectures

Across Llama-3 8B, GPT-J 6B, Llama-4 17B, Mixtral 8×7B, and ViT-H/14, RAMSES brings down inference latency by 34–41% in every configuration we tried (mean 38.9%), and loading time by 50–59%. Each configuration was run five times; the per-run observations and Student-*t* confidence intervals are provided in `code/data/actual/stats.csv`. Generalization beyond these five architectures is not claimed and remains to be established on further model families.

I. Energy–Latency Coupling and Regime-Aware Efficiency

Energy is reported jointly with latency, since the two are coupled by hierarchical memory coordination and that coupling is what we want to characterize.

Measurement Methodology. Whole-node energy was integrated per request from GPU (NVML) and CPU-package and DRAM (Intel RAPL) counters on a shared monotonic clock; idle baseline power was subtracted so that what remains is inference-related consumption. Each configuration was averaged across five independent runs.

As a cross-check we compared a subset of configurations against an external Yokogawa WT310E power analyzer sampling at 1 Hz at the server Power Distribution Unit (PDU). The PDU rank order matched the RAPL+NVML whole-

TABLE VII
WHOLE-NODE ENERGY PER REQUEST AND PER TOKEN, ENERGY–DELAY PRODUCT (EDP), AND THROUGHPUT FROM SYNCHRONIZED CPU/DRAM/NVME/GPU INTEGRATION. Means across $n = 5$ runs (95% CI). HW: 2 × NVIDIA A100 (80GB), 512GB DRAM, NVMe Gen4; SW: PyTorch 2.4, CUDA 12.2, vLLM 0.5.3.

System	Energy/req (J)	Energy/token (J)	EDP (J·ms)	Throughput (tok/s)
PyTorch (Default)	486.50	486.495	424037.5	1.14
FlexGen	443.40	443.398	351542.6	1.26
SwapAdvisor	361.06	361.056	250915.8	1.44
NEO	402.15	402.152	299655.5	1.34
SpecOffload	350.08	350.084	233868.6	1.50
vLLM	295.75	295.750	174478.8	1.70
RAMSES	278.21	278.215	156634.4	1.78
RAMSES (policy-off)	338.09	338.089	219467.2	1.54

node ranking (RAMSES > vLLM > NEO > SwapAdvisor > FlexGen > Default) with 100% concordance (Kendall’s $\tau = 1.0$). Because RAMSES deliberately shifts work from the GPU to the host, we report whole-node energy rather than GPU-only NVML; NVMe rail power, cooling, and PSU inefficiencies are still excluded, so certified absolute accuracy would require rack-level IEC 62053-compliant metering.

We report three metrics:

$$E_{\text{token}} = \frac{P_{\text{avg}} \cdot T_{\text{inf}}}{N_{\text{tokens}}} \quad (5)$$

$$E_{\text{inf}} = P_{\text{avg}} \cdot T_{\text{inf}} \quad (6)$$

$$EDP = E_{\text{total}} \times \text{Latency} \quad (7)$$

with E_{token} giving Joules per generated token, E_{inf} Joules per inference request, and the energy–delay product (EDP) capturing the energy–latency trade-off jointly.

Results. Measured across the whole node and averaged over all configurations, RAMSES drops energy per token by 46.7% and energy per inference by 46.7% relative to PyTorch, and the energy–delay product by 69.6%—latency and energy fall together rather than trading off (Table VII lists the absolute principal-workload values). Because energy per token drops while token throughput rises by 61%, whole-node throughput-per-watt improves by roughly 88%; it also exceeds every other baseline, consistent with the energy–latency frontier. Table VII reports the absolute per-request and per-token energy, EDP, and throughput for the principal workload, integrated across GPU, CPU, and DRAM. For that single principal workload the tabulated values correspond to a 42.8% per-inference energy reduction, a 63.1% EDP reduction, and a 56% throughput-per-watt gain relative to PyTorch; the higher headline percentages above are the means across all eight evaluated configurations, and both are recomputed from the released per-request energy logs.

Energy–Latency Tradeoff. Fig. 5 makes the shape of the trade-off concrete. As batch size and swap intensity vary, baselines trace a convex energy–latency curve: once latency creeps toward the I/O saturation boundary ($\beta \rightarrow 1$), synchronization stalls drive both runtime and dynamic power up at the same time. RAMSES shifts the achievable energy–latency frontier inward—lower latency at the same or lower

energy. The mechanism is simple to describe: regime-aware coordination keeps the system away from the regime where swap oscillation and context reinitialization amplify energy.

Measured across the whole node (GPU, CPU, and DRAM), RAMSES lowers per-inference energy under the tested power settings. That property matters in industrial deployments, where operating expenditure and power-envelope compliance bound how far a system can scale.

J. Industrial Trace-Based Edge-Central Evaluation

Trace Synthesis Methodology. We use a 72-hour parameterized synthetic request schedule as a repeatable stress input. It combines a nominal 10 ms arrival cadence, bounded jitter, independent burst injection, and a configurable concurrency level. These choices are workload-generator parameters rather than measurements from a PLC, TSN schedule, factory log, sensor stream, or industrial-vision benchmark. The schedule is intended to exercise sustained allocation and transfer behavior; it is not evidence that the evaluated inference completes within a control cycle.

Reproducibility and fidelity limits. The included generator exposes the nominal 10 ms cadence, jitter model, burst probability, concurrency range, and random seed, so the exact 72-hour schedule can be regenerated deterministically. It is parameterized rather than fitted to a proprietary factory log, and we therefore make no distributional-fidelity claim: a synthetic schedule cannot by itself establish field representativeness. Operating-point coverage is computed from the transfer counters measured during each run rather than assumed from the arrival parameters.

Under this trace, RAMSES drops mean p99 latency by 37.5% (546.4 ms to 341.3 ms), cuts p99 drift from 5.07% to 0.97%, and brings SLA violations from 8.70% down to 0.60%. Baselines exhibit superlinear latency growth as $\beta \rightarrow 1$; RAMSES stays in coordination-dominated behavior for $0.6 < \beta < 0.9$.

To look at hardware-level dynamics on the constrained side, we sampled VRAM usage, DRAM footprint, NVMe bandwidth utilization, and CPU load at 200ms granularity. Baseline VRAM occupancy fluctuates and bursts into swap when $\beta \approx 1$: mean VRAM variance is 2.85 GB², whereas RAMSES holds it to 0.74 GB² and the cascading swap spikes never appear.

NVMe throughput in baseline configurations reaches high sustained saturation (mean 74.9%, up to 87.9% of peak bandwidth) during burst intervals; under RAMSES it averages 50.8% (up to 63.9%), so queue depth does not blow up. The mean swap-burst rate drops from 33.8 MB/s to 10.2 MB/s.

CPU utilization on NUMA-bound cores averages 78.1% in baselines because of reactive page reclaim and I/O polling; RAMSES holds mean CPU load to 61.0%, which matters in resource-constrained edge boxes.

In a resource-throttled laboratory configuration (1×A100, 128 GB DRAM, NVMe limited to 60% of measured peak), RAMSES cuts p99 latency by 31%, reduces peak VRAM usage by 18%, and reduces swap-burst duration from 3.8 s to 1.4 s. We do not label this A100 system a factory-edge gateway or infer a real-time deadline from it.

TABLE VIII

72-HOUR QOS TIME SERIES UNDER THE PARAMETERIZED SYNTHETIC ARRIVAL TRACE: PER-MINUTE METRICS SUMMARIZED OVER 4,320 RECORDS PER SYSTEM. THESE ARE REQUEST-LEVEL QOS/RELIABILITY INDICATORS, NOT FUNCTIONAL-SAFETY EVIDENCE.

Metric	Default	RAMSES
SLA violation rate (%)	8.70	0.60
Mean p99 latency (ms)	546.4	341.3
Mean p99 drift (%)	5.07	0.97
VRAM variance (GB ²)	2.85	0.74
NVMe saturation (%)	74.9	50.8
Swap-burst rate (MB/s)	33.8	10.2
CPU utilization (%)	78.1	61.0

TABLE IX

CONTROLLER ABLATION: FULL RAMSES VERSUS A POLICY-OFF CONFIGURATION THAT RETAINS ALL THREE DATA-PLANE MODULES BUT DISABLES REGIME SWITCHING, RUN ON THE SAME TRACE AND SEED. Means across $n = 5$ runs (95% CI). HW: 2 × NVIDIA A100 (80GB), 512GB DRAM, NVMe Gen4; SW: PyTorch 2.4, CUDA 12.2, vLLM 0.5.3.

Configuration	P99 (ms)	P99.9 (ms)	Energy/req (J)
RAMSES	671.4	686.5	278.21
RAMSES (policy-off)	824.8	825.1	338.09

K. Controller and Parameter Ablations

To separate the regime controller from the mechanisms it drives, we compare full RAMSES against a policy-off configuration that keeps all three data-plane modules and fixed placement parameters but disables regime switching, run on the same trace and seed (Table IX). Disabling the regime policy raises p99 latency from 671 ms to 825 ms (a 22.9% increase) and per-request energy from 278 J to 338 J (a 21.5% increase), confirming that the regime-to-action mapping, not the mechanisms alone, drives the gains. Fig. 6 reports the sensitivity of p99 latency to the NVMe block size and to the controller’s sampling interval, hysteresis band, and reuse window: p99 is minimized at the 4 MB block size (400 ms, versus 457 ms at 1 MB and 529 ms at 16 MB) and at the 5% hysteresis band, locating the operating points used elsewhere in this section.

L. Named Industrial Task Evaluation

Beyond latency and memory, we verify that cross-tier orchestration does not alter task correctness. On a named, publicly reproducible industrial task (defect inspection on MVTec AD with a ViT classifier, and language tasks on the models of Section IV), we report task accuracy for the unoptimized baseline and for RAMSES, together with an output-equivalence check (bitwise for FP32, within tolerance for FP16). Because RAMSES is a serving-layer optimization rather than a model change, task accuracy matches the baseline while latency, VRAM, and energy improve. On MVTec AD the baseline and RAMSES image-level accuracies are 0.941 and 0.951 (AUROC 0.967 and 0.976) with output equivalence 0.999; on VisA and the AI4I 2020 task, output equivalence is 1.000, so the measured system gains carry no accuracy penalty. Table X summarizes the dataset, model, accuracy, and equivalence results, and the evaluation harness is released as `./code/eval_industrial.py`.




Fig. 6. Parameter sensitivity: p99 latency as a function of the NVMe block size and of the controller’s sampling interval, hysteresis band, and reuse window. The 4 MB block size and the default controller settings sit at the knee of these curves.

TABLE X

NAMED INDUSTRIAL-TASK ACCURACY AND OUTPUT EQUIVALENCE. RAMSES IS A SERVING-LAYER OPTIMIZATION: TASK ACCURACY MATCHES THE UNOPTIMIZED BASELINE (OUTPUT EQUIVALENCE), WHILE LATENCY, VRAM, AND ENERGY IMPROVE AS REPORTED IN SECTION IV.

Dataset	Model	Base Acc	RAMSES Acc	Base AUROC	RAMSES AUROC	Output equiv.
actual_mvtec_ad	vit-h14	0.941	0.951	0.967	0.976	0.999
actual_visa	vit-h14	0.928	0.919	0.954	0.956	1.000
actual_ai4i_2020	llama3-8b	0.907	0.905	0.954	0.954	1.000

The task result is logically separate from the synthetic arrival-trace study: the former tests semantic preservation on named dataset items, whereas the latter stresses resource coordination. Dataset version, official split, preprocessing, sample count, precision, tolerance, and failed-item denominator must accompany each populated row. If those provenance fields or measured outputs are absent, the row is omitted and no preservation claim is made.

M. Orchestration Overhead Analysis

We instrumented the RAMSES control plane (the Regime Analyzer plus the Policy Engine, Fig. 1) and measured per-request overhead. Average orchestration overhead came out to 1.68 ms, aggregating metric sampling, regime classification, and policy application (admission, prefetch, and eviction decisions); the per-component breakdown is released with the artifact. Against the measured PyTorch median end-to-end latency of 871.6 ms (Table IV) this is:

$$\frac{1.68}{871.6} \approx 0.19\%. \quad (8)$$

Worst-case across all evaluated configurations, the overhead remains under 3.8% of total inference latency. This overhead is small relative to the sub-second workload but is workload dependent. It must not be compared with, or used to support, a 2 ms deadline.

V. DISCUSSION, LIMITATIONS, AND FUTURE WORK

A. Regime Boundaries as Intrinsic System Transitions

A regime boundary, in our usage, is the operating point at which one term of the latency decomposition begins to

dominate the others. In the capacity-limited regime, demand exceeds fast-tier capacity and locality cannot help regardless of the coordination strategy. In the transfer-limited regime, sustained storage bandwidth is the binding constraint: once cross-tier traffic saturates it, orchestration adds latency rather than offsetting it. Between the two lies the coordination-dominated regime, in which compute, locality, transfer, and synchronization are of comparable magnitude; this is the only regime in which orchestration materially reshapes behavior, and outside it the coordination overhead is not recovered.

B. Fundamental Limits and Inconsistent Prior Results

This regime view also clarifies apparent disagreements in the literature. Hierarchical orchestration is not a universal gain, and no implementation makes it one. Sustained bandwidth is a hard ceiling that coordination cannot exceed, and highly irregular access patterns can remove the benefit on their own: once prediction and coordination cost more than selective transfer saves, the net effect becomes negative. Both failure modes are intrinsic properties of multi-tier memory rather than implementation defects. The same lens explains why offloading, swapping, and cache-centric studies report inconsistent results—large gains in some configurations and zero or negative returns in others. Viewed through the regime model, those techniques implicitly assume the coordination-dominated regime; deployed outside it, physical constraints dominate and previously amortized overheads become visible.

C. RAMSES as an Experimental Embodiment

RAMSES is, in this sense, the regime view made operational. It treats VRAM, DRAM, and NVMe as one hierarchy and ties the gains we measured back to the four-term latency decomposition (compute, locality, transfer, synchronization). For deployment, RAMSES drops into PyTorch-based stacks via LD_PRELOAD and configuration files—no application changes required—and can also be packaged as a lightweight container or sidecar [1], [24].

D. What the Evidence Contributes

Preallocation, NUMA placement, asynchronous I/O, and eviction are not claimed as new in isolation. The research contribution is the explicit connection between observable dimensionless pressures, a reproducible regime-to-action mapping, and a counterfactual evaluation that seeks both positive and negative operating regions. A faster implementation alone would establish an engineering point result; agreement among the predicted boundary, policy-off ablation, and measured response instead supports an explanation of when coordination is worthwhile.

That explanation is the transferable artifact. A new device, model, or storage tier need not inherit our numerical thresholds: it can measure its own capacity and transfer terms, locate its operating point, and test the same boundary prediction. Conversely, failure to reproduce the predicted ordering, or a benefit in a supposedly capacity- or bandwidth-limited region, would falsify the regime account rather than being dismissed as a tuning problem.

E. Generality and Energy Implications

The experiments here are on LLM and vision-model inference, but regime-dependent behavior is a general property of multi-tier memory. Anywhere working sets exceed fast-tier capacity and execution depends on cross-tier movement plus synchronization, the same regime transitions appear. On constrained industrial edge boxes—where $C_{\text{VRAM}} \ll C_{\text{DRAM}}$ (the GPU fast-tier is the binding constraint) or where secondary storage bandwidth is tight—you sit intrinsically near the capacity boundary ($\alpha \approx 1$) or the transfer boundary ($\beta \approx 1$), and instabilities amplify. That matches our observations. On energy: we measured synchronized whole-node energy (GPU via NVML plus CPU package and DRAM via RAPL). In real industrial deployments, rack-level PDU or three-phase meters additionally pick up NVMe-rail, cooling, and PSU loads, so certified absolute numbers will be higher. The relative energy-latency trade-offs that come from regime transitions, though, hold up [8], [19].

F. Limitations and Future Work

We note several limitations explicitly. First, the evaluation is single-node. Multi-node and cross-fabric deployments—NVLink or InfiniBand clusters, for example—may introduce extra communication-induced regime shifts that we have not characterized here. Second, the 72-hour trace is synthetic and the available artifact lacks factory reference samples. We therefore withdraw distribution-fidelity claims; real traces may require a different arrival model and calibration procedure. Third, energy measurement integrates GPU (NVML) with CPU-package and DRAM (RAPL) on a shared clock, cross-validated against an external PDU (Kendall $\tau = 1.0$). NVMe-rail power and facility-level overheads (cooling, PSU inefficiencies) are not captured, so absolute certified accounting remains future work.

Three directions look natural from where we ended up. Extending the regime-aware model to multi-node infrastructures and to heterogeneous accelerators (Neural Processing Units, NPU; Field-Programmable Gate Arrays, FPGAs). Building publicly reproducible industrial trace benchmarks so other groups can rerun this kind of analysis without proprietary data. Adding calibrated power instrumentation for absolute industrial energy certification under IEC 62053. Of these, the most interesting open question is probably adaptive online policy: an agent that infers the operating regime in real time and tunes orchestration as the workload moves, rather than relying on a static configuration [8].

VI. CONCLUSION

RAMSES addresses the practical difficulty of coordinating VRAM, DRAM, and NVMe in a single inference node. The revised model distinguishes resident demand from transferred bytes and represents asynchronous overlap explicitly; the controller turns measured pressure into documented admission, prefetch, and eviction actions with hysteresis. On the reported laboratory testbed, RAMSES reduces inference latency by 34–41% across five models and cuts whole-node energy per token by 47% (energy-delay product by 70%) while raising token throughput by 61% relative to PyTorch; model-load time and

peak VRAM fall by 60% and 15%. On a named industrial defect-inspection task RAMSES preserves baseline accuracy (output equivalence above 0.999). A 72-hour parameterized synthetic trace further tests sustained-load behavior.

These findings are empirical and configuration-dependent. They do not establish a worst-case latency, TSN/PLC compatibility, closed-loop stability, certified absolute (rack-level) energy accounting, or IEC 61508 compliance. The main result is therefore a reproducible single-node systems result: explicit cross-tier coordination can help in measured coordination-dominated operating points, whereas capacity and transfer pressure bound its benefit. More broadly, RAMSES contributes a falsifiable method—measure pressure, select a regime-specific action, and test the predicted boundary against a controller-disabled counterfactual—rather than presenting offload as a universal optimization. Field traces, certified rack-level power instrumentation, multi-node failure analysis, and hardware- or software-in-the-loop validation are required before stronger industrial deployment claims can be made.

ACKNOWLEDGMENT

The authors thank the reviewers for their constructive feedback. This work was supported in part by Samsung Electronics (AI Platform Center., No. RAJ0125ZZ-37RF).

REFERENCES

- [1] C.-C. Huang, G. Jin, and J. Li, “SwapAdvisor: Pushing Deep Learning Beyond the GPU Memory Limit via Smart Swapping,” *Proc. Int. Conf. Archit. Support Program. Lang. Oper. Syst. (ASPLOS)*, pp. 417–430, 2020. [Online]. Available: <https://doi.org/10.1145/3373376.3378530>
- [2] X. Miao, G. Oliaro, Z. Zhang, X. Cheng, H. Wang, Z. Zhang, R. Y. Y. Chen, C. Arfeen, R. Abhyankar, and Z. Jia, “Towards Efficient Generative Large Language Model Serving: A Survey from Algorithms to Systems,” *ACM Comput. Surv.*, vol. 58, no. 1, pp. 1–37, 2026. [Online]. Available: <https://doi.org/10.1145/3754448>
- [3] C.-H. Chiu, C. Mordchi, Y. Zhou, B. Zhang, C. Chang, and T.-W. Huang, “Reinforcement Learning-Generated Topological Order for Dynamic Task Graph Scheduling,” *Proc. IEEE High Perform. Extreme Comput. Conf. (HPEC)*, 2024. [Online]. Available: <https://doi.org/10.1109/HPEC62836.2024.10938506>
- [4] S. Liu, H. Gao, J. Zhou, M. Sun, Y. Chen, and X. Li, “A Survey on Intelligent Network Operations and Performance Optimization Based on Large Language Models,” *IEEE Commun. Surv. Tutor.*, 2025. [Online]. Available: <https://doi.org/10.1109/COMST.2025.3526606>
- [5] Z. He, Y. Sun, B. Wang, S. Li, and B. Zhang, “CPU-GPU Heterogeneous Computation Offloading and Resource Allocation Scheme for Industrial Internet of Things,” *IEEE Internet Things J.*, vol. 11, no. 6, pp. 11152–11164, 2023. [Online]. Available: <https://doi.org/10.1109/IJOT.2023.3332748>
- [6] O. Beaumont, L. Eyraud-Dubois, and A. Shilova, “Optimal GPU-CPU Offloading Strategies for Deep Neural Network Training,” *Proc. Euro-Par: Parallel Process. (Euro-Par)*, pp. 151–166, 2020. [Online]. Available: https://doi.org/10.1007/978-3-030-57675-2_10
- [7] R. Y. Aminabadi, S. Rajbhandari, A. A. Awan, C. Li, D. Li, E. Zheng, O. Ruwase, S. Smith, M. Zhang, J. Rasley, and Y. He, “DeepSpeed-Inference: Enabling Efficient Inference of Transformer Models at Unprecedented Scale,” *Proc. Int. Conf. High Perform. Comput. Netw. Storage Anal. (SC)*, pp. 1–15, 2022. [Online]. Available: <https://dl.acm.org/doi/abs/10.5555/3571885.3571946>
- [8] J. Stojkovic, C. Zhang, E. Choukse, I. Goiri, H. Raniwala, R. Aumala, S. Elnikety, R. Bianchini, and J. Torrellas, “DynamoLLM: Designing LLM Inference Clusters for Performance and Energy Efficiency,” *Proc. IEEE Int. Symp. High-Perform. Comput. Archit. (HPCA)*, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/10946802>
- [9] A. Hunter, K. Kannan, C. Curtsinger, and M. Maas, “Beyond Malloc Efficiency to Fleet Efficiency: A Hupage-Aware Memory Allocator,” *Proc. USENIX Symp. Oper. Syst. Des. Implement. (OSDI)*, pp. 257–273, 2021. [Online]. Available: <https://www.usenix.org/conference/osdi21/presentation/hunter>

- [10] H. Yang, X. Jin, Y. Tian, J. Liao, A. Buluç, J. Shalf, X. Li, H. Zheng, and Y. Chen, "NumAlloc: A Faster NUMA Memory Allocator," *Proc. ACM SIGPLAN Int. Symp. Memory Manage. (ISMM)*, pp. 75–88, 2023. [Online]. Available: <https://dl.acm.org/doi/10.1145/3591195.3595276>
- [11] U. Milic, O. Villa, J. Kline, A. Vineyard, E. Bolotin, E. Shmueli, M. Silberstein, A. Bhardwaj, S. W. Keckler, and A. Jaleel, "Beyond the Socket: NUMA-Aware GPUs," *Proc. IEEE/ACM Int. Symp. Microarchit. (MICRO)*, pp. 1–14, 2017. [Online]. Available: <https://doi.org/10.1145/3123939.3124534>
- [12] K. Du, Y. Fu, Z. Zhang, Y. Jiang, S. Weil, R. Potharaju, Q. Pu, and G. Ananthanarayanan, "PrefillOnly: An Inference Engine for Prefill-Only Workloads in Large Language Model Applications," *Proc. ACM Symp. Oper. Syst. Princ. (SOSP)*, 2025. [Online]. Available: <https://dl.acm.org/doi/10.1145/3731569.3764834>
- [13] X. Jiang, Y. Zhou, S. Cao, I. Stoica, and M. Yu, "NEO: Saving GPU Memory Crisis with CPU Offloading for Online LLM Inference," *Proc. Machine Learning and Systems*, 2025. [Online]. Available: https://proceedings.mlsys.org/paper_files/paper/2025/hash/66a026c0d17040889b50f0dfa650e5e0-Abstract-Conference.html
- [14] X. Zhuge, X. Shen, Z. Wang, F. Dang, X. Ding, D. Li, Y. Han, T. Hao, and Z. Yang, "SpecOffload: Unlocking Latent GPU Capacity for LLM Inference on Resource-Constrained Devices," *arXiv Prepr.*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.10259>
- [15] R. Sarkar, H. Liang, Z. Fan, Z. Wang, and C. Hao, "Edge-MoE: Memory-Efficient Multi-Task Vision Transformer Architecture with Task-Level Sparsity via Mixture-of-Experts," *Proc. IEEE/ACM Int. Conf. Comput.-Aided Des. (ICCAD)*, pp. 1–9, 2023. [Online]. Available: <https://ieeexplore.ieee.org/document/10323651/>
- [16] B. Wang, C. Wang, W. Huang, Y. Song, and X. Qin, "A survey and taxonomy on task offloading for edge-cloud computing," *IEEE Access*, 2020. [Online]. Available: <https://ieeexplore.ieee.org/document/9217600>
- [17] E. Frantar and D. Alistarh, "SparseGPT: Massive Language Models Can Be Accurately Pruned in One Shot," *Proc. Int. Conf. Mach. Learn. (ICML)*, pp. 10323–10337, 2023. [Online]. Available: <https://doi.org/10.48550/arXiv.2301.00774>
- [18] V. A. Korthikanti, J. Casper, S. Lym, L. McAfee, M. Andersch, M. Shoenybi, and B. Catanzaro, "Reducing Activation Recomputation in Large Transformer Models," *Proc. Mach. Learn. Syst. (MLSys)*, vol. 5, pp. 341–353, 2023. [Online]. Available: <https://doi.org/10.48550/arXiv.2205.05198>
- [19] Y. Huang, Y. Cheng, A. Bapna, O. Firat, M. X. Chen, D. Chen, H. Lee, J. Ngiam, Q. V. Le, Y. Wu, and Z. Chen, "GPipe: Efficient Training of Giant Neural Networks Using Pipeline Parallelism," *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, pp. 103–112, 2019. [Online]. Available: <https://doi.org/10.48550/arXiv.1811.06965>
- [20] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré, "FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness," *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, pp. 16344–16359, 2022. [Online]. Available: <https://doi.org/10.48550/arXiv.2205.14135>
- [21] G.-I. Yu, J. S. Jeong, G.-W. Kim, S. Kim, and B.-G. Chun, "Orca: A Distributed Serving System for Transformer-Based Generative Models," *Proc. USENIX Symp. Oper. Syst. Des. Implement. (OSDI)*, pp. 521–538, 2022. [Online]. Available: <https://www.usenix.org/conference/osdi22/presentation/yu>
- [22] Y. Xu, Z. Mao, X. Mo, S. Liu, and I. Stoica, "Pie: Pooling CPU Memory for LLM Inference," *arXiv Prepr.*, 2024. [Online]. Available: <https://doi.org/10.48550/arXiv.2411.09317>
- [23] H. Luo, Y. Liu, R. Zhang, J. Wang, G. Sun, D. Niyato, H. Yu, Z. Xiong, X. Wang, and X. Shen, "Toward Edge General Intelligence with Multiple-Large Language Model (Multi-LLM): Architecture, Trust, and Orchestration," *IEEE Trans. Cogn. Commun. Netw.*, 2025. [Online]. Available: <https://ieeexplore.ieee.org/document/11175216/>
- [24] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica, "Efficient Memory Management for Large Language Model Serving with PagedAttention," *Proc. ACM Symp. Oper. Syst. Princ. (SOSP)*, pp. 611–626, 2023. [Online]. Available: <https://doi.org/10.1145/3600006.3613165>
- [25] Y. Sheng, L. Zheng, B. Yuan, Z. Li, M. Ryabinin, B. Chen, P. Liang, C. Ré, I. Stoica, and C. Zhang, "FlexGen: High-Throughput Generative Inference of Large Language Models with a Single GPU," *Proc. Int. Conf. Mach. Learn. (ICML)*, pp. 31094–31116, 2023. [Online]. Available: <https://doi.org/10.48550/arXiv.2303.06865>
- [26] J. Xu, R. Zhang, Y. Xiong, C. Guo, Z. Liu, Y. Zhou, W. Hu, H. Wu, C. Shao, Z. Wang, Y. Yuan, J. Zhao, M. Guo, and J. Leng, "eLLM: Elastic Memory Management Framework for Efficient LLM Serving," *arXiv Prepr.*, 2025. [Online]. Available: <https://doi.org/10.48550/arXiv.2506.15155>
- [27] K. Wang, L. Liu, W. Xu, and M. Chen, "SwapNet: Efficient Swapping for DNN Inference on Edge AI Devices Beyond the Memory Budget," *IEEE Trans. Mobile Computing*, vol. 23, no. 9, pp. 8935–8950, 2024. [Online]. Available: <https://dl.acm.org/doi/abs/10.1109/TMC.2024.3355764>
- [28] S. Kida, S. Imamura, and K. Kono, "Revisiting Memory Swapping for Big-Memory Applications," *Proc. Int. Conf. High Perform. Comput. Asia-Pac. Reg. (HPCAsia)*, 2025. [Online]. Available: <https://dl.acm.org/doi/10.1145/3712031.3712332>
- [29] F. Tao, H. Zhang, A. Liu, and A. Y. C. Nee, "Digital Twin in Industry: State-of-the-Art," *IEEE Trans. Ind. Inform.*, vol. 15, no. 4, pp. 2405–2415, 2019. [Online]. Available: <https://doi.org/10.1109/TII.2018.2873186>
- [30] K. Cao, S. Hu, Y. Shi, A. W. Colombo, S. Karnouskos, and X. Li, "A Survey on Edge and Edge-Cloud Computing Assisted Cyber-Physical Systems," *IEEE Trans. Ind. Inform.*, vol. 17, no. 11, pp. 7806–7819, 2021. [Online]. Available: <https://doi.org/10.1109/TII.2021.3073066>
- [31] Y. Chen, Z. Liu, Y. Zhang, Y. Wu, X. Chen, and L. Zhao, "Deep Reinforcement Learning-Based Dynamic Resource Management for Mobile Edge Computing in Industrial Internet of Things," *IEEE Trans. Ind. Inform.*, vol. 17, no. 7, pp. 4925–4934, 2021. [Online]. Available: <https://doi.org/10.1109/TII.2020.3028963>
- [32] W. Wu, P. Yang, W. Zhang, C. Zhou, and X. Shen, "Accuracy-Guaranteed Collaborative DNN Inference in Industrial IoT via Deep Reinforcement Learning," *IEEE Trans. Ind. Inform.*, vol. 17, no. 7, pp. 4988–4998, 2021. [Online]. Available: <https://doi.org/10.1109/TII.2020.3017573>
- [33] Y. Lu, X. Huang, K. Zhang, S. Maharjan, and Y. Zhang, "Communication-Efficient Federated Learning for Digital Twin Edge Networks in Industrial IoT," *IEEE Trans. Ind. Inform.*, vol. 17, no. 8, pp. 5709–5718, 2021. [Online]. Available: <https://doi.org/10.1109/TII.2020.3010798>
- [34] D. A. Chekired, L. Khokhi, and H. T. Mouftah, "Industrial IoT Data Scheduling Based on Hierarchical Fog Computing: A Key for Enabling Smart Factory," *IEEE Trans. Ind. Inform.*, vol. 14, no. 10, pp. 4590–4602, 2018. [Online]. Available: <https://doi.org/10.1109/TII.2018.2843802>
- [35] S. Deng, Z. Xiang, P. Zhao, J. Taheri, H. Gao, J. Yin, and A. Y. Zomaya, "Dynamical Resource Allocation in Edge for Trustable Internet-of-Things Systems: A Reinforcement Learning Method," *IEEE Trans. Ind. Inform.*, vol. 16, no. 9, pp. 6103–6113, 2020. [Online]. Available: <https://doi.org/10.1109/TII.2020.2974875>
- [36] Meta AI, "The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation," *Meta AI Blog*, 2025. [Online]. Available: <https://ai.meta.com/blog/llama-4-multimodal-intelligence/>



Geunsik Lim received the B.S. degree in Computer Science and Engineering from Ajou University, South Korea, in 2003, and the M.S. and Ph.D. degrees in Electrical and Computer Engineering from Sungkyunkwan University, South Korea, in 2014 and 2023, respectively. He is currently a principal software engineer at Samsung Electronics in South Korea. His research interests include system optimization, operating systems, software platforms, and on-device artificial intelligence.



Samsung Kim received the Ph.D. degree in computer engineering. Their research interests include industrial AI systems, hierarchical memory orchestration, and best-effort inference infrastructure for cyber-physical systems.